

Module	4G10	Title of report	4G10 Coursework 2
Date submitted: 25 Dec		Assessment for this module is <input checked="" type="checkbox"/> 100% / <input type="checkbox"/> 25% coursework of which this assignment forms 50 %	
UNDERGRADUATE and POST GRADUATE STUDENTS			
Candidate number:	5641G		<input checked="" type="checkbox"/> Undergraduate <input type="checkbox"/> Post graduate

Feedback to the student		Very good	Good	Needs improvmt
<input type="checkbox"/> See also comments in the text				
C O N T E N T	<b>Completeness, quantity of content:</b> Has the report covered all aspects of the lab? Has the analysis been carried out thoroughly?			
	<b>Correctness, quality of content</b> Is the data correct? Is the analysis of the data correct? Are the conclusions correct?			
	<b>Depth of understanding, quality of discussion</b> Does the report show a good technical understanding? Have all the relevant conclusions been drawn?			
	Comments:			
P R E S E N T A T I O N	<b>Attention to detail, typesetting and typographical errors</b> Is the report free of typographical errors? Are the figures/tables/references presented professionally?			
	Comments:			

# 4G10 Coursework 2

Author: 5641G

## 1 Introduction

This report addresses the task of predicting hand kinematics from neural data recorded in the primary motor cortex (M1) of a macaque monkey. The dataset consists of neural spike counts from 162 cortical neurons and corresponding 2D hand velocity measurements across multiple reaching trials. Each trial begins at a 'go' cue and lasts 800 ms, divided into 16 bins of 50 ms. The data is partitioned into training and testing sets, with 400 and 100 trials respectively. Various decoding techniques are explored, including Gaussian smoothing combined with ridge regression and Kalman filtering. The results are analyzed to assess prediction accuracy and computational feasibility for online decoding applications.

## 2 Baseline Decoder: Simple Gaussian Smoothing + Linear Regression

The baseline decoding approach consists of two stages: smoothing the spike count time series using a Gaussian filter and applying ridge regression for linear decoding. Ridge regression minimizes the following cost function to balance data fidelity and regularization:

$$\mathcal{L} = \|\mathbf{V} - \mathbf{W}\tilde{\mathbf{X}}\|_2^2 + \lambda\|\mathbf{W}\|_2^2, \quad (1)$$

where  $\mathbf{V}$  is the  $2 \times (400 \cdot 16)$  matrix of hand velocities,  $\tilde{\mathbf{X}}$  is the  $N \times (400 \cdot 16)$  matrix of smoothed neural spike counts,  $\lambda$  is the regularization parameter, and  $\mathbf{W}$  represents the decoder weights. The closed-form solution for  $\mathbf{W}$  is given by:

$$\mathbf{W}^* = \mathbf{V}\tilde{\mathbf{X}}^\top (\tilde{\mathbf{X}}\tilde{\mathbf{X}}^\top + \lambda\mathbf{I}_N)^{-1}. \quad (2)$$

Here,  $\tilde{\mathbf{X}}\tilde{\mathbf{X}}^\top$  forms a Gram matrix, as it is not mean-centered.

### 2.1 Choice of $\lambda$

To determine an optimal value for  $\lambda$ , several values were tested with  $\sigma$  fixed at 40 ms. Table 1 summarizes the results. The optimal value,  $\lambda = 75$ , balances regularization and prediction accuracy:

$\lambda$	0	0.1	0.5	1	2	10	50	100	200	$10^4$
$R^2$	0.422	0.422	0.422	0.423	0.423	0.425	0.428	0.428	0.426	0.390

Table 1: Selection of  $\lambda$  with  $\sigma = 40$  ms

### 2.2 Effect of Smoothing Window Length $\sigma$

The impact of smoothing window length  $\sigma$  on decoding accuracy was analyzed using offline ridge regression with  $\lambda = 75$ . As shown in Table 2 and the blue curve in Figure 1, the quality of hand velocity predictions peaks at  $\sigma \approx 60$  ms. Shorter windows (e.g.,  $\sigma = 20$  ms) fail to capture sufficient temporal structure, likely overfitting noise. Conversely, overly large windows (e.g.,  $\sigma = 80$  ms) blur meaningful temporal variations, degrading performance.

$\sigma$ (ms)	20	25	30	40	50	55	60	65	70	75	80
$R^2$	0.345	0.375	0.401	0.428	0.441	0.444	0.445	0.445	0.443	0.440	0.437

Table 2: Effect of  $\sigma$  on offline decoding performance with  $\lambda = 75$

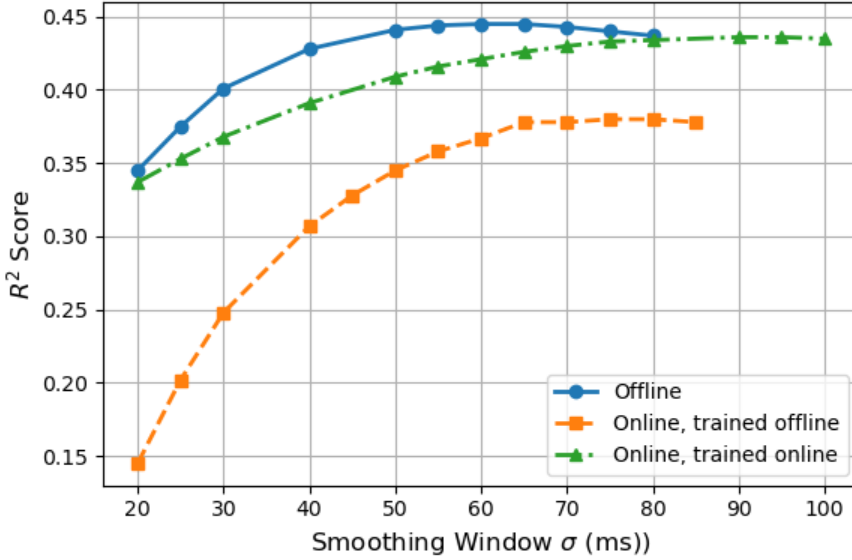


Figure 1: Effect of  $\sigma$  on  $R^2$  for baseline decoder and its variants.

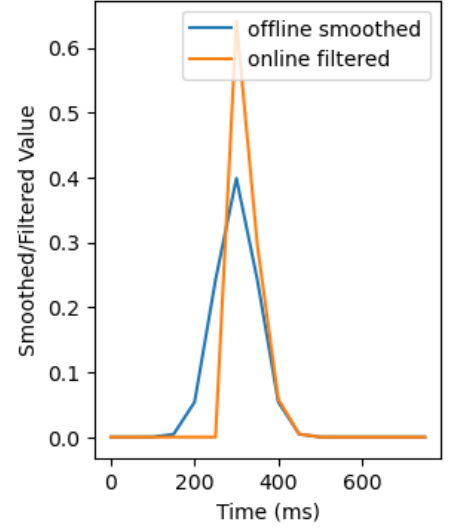


Figure 2: Comparison of online filtering and offline smoothing.

### 2.3 Suitability for Online Decoding

The simple decoding strategy is computationally efficient and feasible for online decoding in a brain-machine interface (BMI) context. However, its accuracy drops significantly when applied online, as shown in Table 3 and the orange curve in Figure 1. This decline arises because online filtering does not access future information, leading to suboptimal smoothing. Figure 2 illustrates this limitation, where online filtering differs markedly from offline smoothing for a single spike.

$\sigma$ (ms)	20	30	40	50	60	65	70	75	80	85
$R^2$	0.145	0.248	0.307	0.345	0.367	0.378	0.378	0.380	0.380	0.378

Table 3: Online decoding performance using offline-trained weights with  $\lambda = 75$

Additionally, neural tuning properties may drift over time, reducing decoding accuracy. Adaptive updates or retraining of the model may mitigate this issue.

### 2.4 Improving Online Decoding

A simple improvement involves training the weights  $\mathbf{W}$  on online-filtered training data instead of smoothed data. This modification aligns the training process with the online setting, resulting in significant performance improvements. As shown in Table 4 and green curve in Figure 1, the  $R^2$  score improves to 0.436 at  $\sigma \approx 90$  ms, approaching the optimal offline performance of 0.445.

$\sigma$ (ms)	20	30	40	50	60	70	80	90	95	100
$R^2$	0.337	0.368	0.391	0.409	0.421	0.430	0.434	0.436	0.436	0.435

Table 4: Online decoding performance using online-trained weights with  $\lambda = 75$

### 2.5 Temporal Shift in Hand Velocity Data

The hand velocity data in `data['hand_train']` was intentionally shifted backward by 120 ms relative to the neural data. This adjustment likely accounts for the natural delay in the motor control process. Neural activity in the primary motor cortex (M1) precedes the execution of motor actions because it encodes planning and initiation signals for movement. By aligning the neural data with the hand velocity data in this manner, the dataset ensures that neural signals related to planning and initiation are temporally synchronized with the corresponding behavioral outputs. Without this shift, the behavioral data would lag behind the neural data, introducing a temporal mismatch that could compromise the performance of decoding algorithms.

### 3 Kalman Filter-Based Decoding

We now introduce a more advanced decoding method based on a Kalman filter and smoother.

#### 3.1 An Autoregressive Prior for Hand Kinematics

A pre-trained 10-dimensional linear latent dynamical system (LDS) is utilized to model hand velocity data from the training set. The generative model is defined as follows:

$$\mathbf{z}_{k,0} \sim \mathcal{N}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0) \quad (\text{initial prior}) \quad (3)$$

$$\mathbf{z}_{k,t+1} = \mathbf{A}\mathbf{z}_{k,t} + \boldsymbol{\epsilon}_{k,t+1}, \quad \boldsymbol{\epsilon}_{k,t+1} \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}) \quad (\text{state transition}) \quad (4)$$

$$\mathbf{v}_{k,t} = \mathbf{C}\mathbf{z}_{k,t} + \boldsymbol{\eta}_{k,t}, \quad \boldsymbol{\eta}_{k,t} \sim \mathcal{N}(\mathbf{0}, \mathbf{R}) \quad (\text{observation model}) \quad (5)$$

Here,  $\mathbf{z}_{k,t} \in \mathbb{R}^{10}$  represents the latent state at time bin  $t$  of trial  $k$ , and  $\mathbf{v}_{k,t} \in \mathbb{R}^2$  represents the corresponding hand velocity.

The Kalman smoother combines forward (filtering) and backward (smoothing) passes to estimate latent states. The goal is to implement a Kalman smoother to compute the mean  $\hat{z}_{k,1:T}$  of the smoothing distribution  $p(z_{k,t}|v_{k,1:T})$  for each trial  $k$  in the training set.

##### 3.1.1 Kalman Filter (Forward Pass)

The Kalman filter recursively estimates the latent state  $z_t$  based on observations up to time  $t$ :

$$\mathbf{P}_t = \mathbf{A}\boldsymbol{\Sigma}_t\mathbf{A}^\top + \mathbf{Q}, \quad (6)$$

$$\mathbf{K}_t = \mathbf{P}_t\mathbf{C}^\top (\mathbf{R} + \mathbf{C}\mathbf{P}_t\mathbf{C}^\top)^{-1}, \quad (7)$$

$$\boldsymbol{\Sigma}_{t+1} = (\mathbf{I} - \mathbf{K}_t\mathbf{C})\mathbf{P}_t, \quad (8)$$

$$\boldsymbol{\mu}_{t+1} = \mathbf{A}\boldsymbol{\mu}_t + \mathbf{K}_t(\mathbf{x}_{t+1} - \mathbf{C}\mathbf{A}\boldsymbol{\mu}_t). \quad (9)$$

Here,  $\mathbf{K}_t$  is the Kalman gain, which balances the prior and the observation likelihood. Observations  $\mathbf{x}_t$  are the hand velocity measurements  $v_{k,t}$ .

##### 3.1.2 Kalman Smoother (Backward Pass)

The Kalman smoother refines the filtered estimates by incorporating future observations:

$$\mathbf{P}_t = \mathbf{A}\boldsymbol{\Sigma}_t\mathbf{A}^\top + \mathbf{Q}, \quad (10)$$

$$\mathbf{G}_t = \boldsymbol{\Sigma}_t\mathbf{A}^\top\mathbf{P}_t^{-1}, \quad (11)$$

$$\tilde{\boldsymbol{\mu}}_t = \boldsymbol{\mu}_t + \mathbf{G}_t(\tilde{\boldsymbol{\mu}}_{t+1} - \mathbf{A}\boldsymbol{\mu}_t), \quad (12)$$

$$\tilde{\boldsymbol{\Sigma}}_t = \boldsymbol{\Sigma}_t + \mathbf{G}_t(\tilde{\boldsymbol{\Sigma}}_{t+1} - \mathbf{P}_t)\mathbf{G}_t^\top. \quad (13)$$

Here,  $\tilde{\boldsymbol{\mu}}_t$  and  $\tilde{\boldsymbol{\Sigma}}_t$  represent the smoothed mean and covariance, respectively.

These updates were implemented and applied to compute smoothed latent state estimates for the training data. The resulting  $\hat{z}_{k,1:T}$  values correspond to the smoothed means  $\tilde{\boldsymbol{\mu}}_t$ , which represent the posterior mean of the latent state distribution after incorporating both past and future observations. These values enables the predictive decoding of hand velocity from neural data.

#### 3.2 Building an LDS Model of Neural Data Using Supervised Learning

In this section, we extend the generative LDS model to neural data:

$$\mathbf{x}_{k,t} = \mathbf{D}\mathbf{z}_{k,t} + \boldsymbol{\xi}_{k,t}, \quad \boldsymbol{\xi}_{k,t} \sim \mathcal{N}(\mathbf{0}, \mathbf{S}) \quad (14)$$

where  $\mathbf{x}_{k,t}$  represents the neural activity (spike counts) for trial  $k$  at time  $t$ , and  $\mathbf{z}_{k,t}$  is the latent state variable, which can be obtained from the Kalman smoother in Section 3.1.

### 3.2.1 Centering the Neural Data

To ensure proper alignment between neural data and latent state dynamics, we first center the neural data:

- Compute the mean activity across trials and time for each neuron in the training set.
- Subtract this mean from both the training and testing neural data.

### 3.2.2 Fitting Parameters

The likelihood parameters  $\mathbf{D}$  (observation matrix) and  $\mathbf{S}$  (observation noise covariance) are estimated by maximizing the average joint log-likelihood:

$$\log p(\hat{\mathbf{z}}_{k,1:T}, \mathbf{x}_{k,1:T}) = \frac{1}{K} \sum_k \log p(\mathbf{x}_{k,1:T} | \hat{\mathbf{z}}_{k,1:T}) + \log p(\hat{\mathbf{z}}_{k,1:T}) \quad (15)$$

$$= \frac{1}{K} \sum_k \left[ \sum_{t=1}^T \log p(\mathbf{x}_{k,t} | \hat{\mathbf{z}}_{k,t}) + \log p(\hat{\mathbf{z}}_{k,1:T}) \right]. \quad (16)$$

where  $K$  is the number of trials,  $\hat{\mathbf{z}}_{k,1:T}$  is the posterior mean from Section 3.1, and  $\mathbf{x}_{k,1:T}$  is the neural data. Since  $p(\hat{\mathbf{z}}_{k,1:T})$  is independent of  $\mathbf{D}, \mathbf{S}$ , we only maximize:

$$\log p(\mathbf{x}_{k,t} | \hat{\mathbf{z}}_{k,t}) = \frac{1}{K} \sum_{k,t} \left[ -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log |\mathbf{S}| - \frac{1}{2} (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t})^\top \mathbf{S}^{-1} (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t}) \right]. \quad (17)$$

#### Gradient w.r.t $\mathbf{D}$

$$\frac{\partial}{\partial \mathbf{D}} \log p = \frac{1}{K} \sum_{k,t} \frac{\partial}{\partial \mathbf{D}} \left[ -\frac{1}{2} (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t})^\top \mathbf{S}^{-1} (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t}) \right] \quad (18)$$

$$= \frac{1}{K} \sum_{k,t} \left[ \mathbf{S}^{-1} \mathbf{x}_{k,t} \hat{\mathbf{z}}_{k,t}^\top - \mathbf{S}^{-1} \mathbf{D} \hat{\mathbf{z}}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right]. \quad (19)$$

Setting this gradient to zero and solving for  $\mathbf{D}^*$ :

$$\mathbf{S}^{-1} \frac{1}{K} \sum_{k,t} \left[ \mathbf{x}_{k,t} \hat{\mathbf{z}}_{k,t}^\top - \mathbf{D}^* \hat{\mathbf{z}}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right] = \mathbf{0} \quad (20)$$

$$\mathbf{D}^* = \left( \sum_{k,t} \mathbf{x}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right) \left( \sum_{k,t} \hat{\mathbf{z}}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right)^{-1}. \quad (21)$$

#### Gradient w.r.t $\mathbf{S}$

Applying the formula  $\frac{\partial}{\partial \mathbf{A}} (\mathbf{x}^\top \mathbf{A}^{-1} \mathbf{x}) = -\mathbf{A}^{-1} \mathbf{x} \mathbf{x}^\top \mathbf{A}^{-1}$ .

$$\frac{\partial}{\partial \mathbf{S}} \log p = \frac{1}{K} \sum_{k,t} \left[ -\frac{1}{2} \mathbf{S}^{-1} + \frac{1}{2} \mathbf{S}^{-1} (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t}) (\mathbf{x}_{k,t} - \mathbf{D} \hat{\mathbf{z}}_{k,t})^\top \mathbf{S}^{-1} \right]. \quad (22)$$

Setting this gradient to zero, solving for  $\mathbf{S}^*$

$$\mathbf{S}^* = \frac{1}{KT} \sum_{k,t} (\mathbf{x}_{k,t} - \mathbf{D}^* \hat{\mathbf{z}}_{k,t}) (\mathbf{x}_{k,t} - \mathbf{D}^* \hat{\mathbf{z}}_{k,t})^\top \quad (23)$$

Substituting  $\mathbf{D}^*$

$$\mathbf{S}^* = \frac{1}{KT} \sum_{k,t} \left[ \mathbf{x}_{k,t} \mathbf{x}_{k,t}^\top - \mathbf{D}^* \hat{\mathbf{z}}_{k,t} \mathbf{x}_{k,t}^\top - \mathbf{x}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \mathbf{D}^{*\top} + \mathbf{D}^* \hat{\mathbf{z}}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \mathbf{D}^{*\top} \right] \quad (24)$$

$$= \frac{1}{KT} \left( \sum_{k,t} \mathbf{x}_{k,t} \mathbf{x}_{k,t}^\top - \mathbf{D}^* \sum_{k,t} \hat{\mathbf{z}}_{k,t} \mathbf{x}_{k,t}^\top \right) \quad (25)$$

### 3.2.3 Final Parameter Estimates

The fitted parameters are:

$$\mathbf{D}^* = \left( \sum_{k,t} \mathbf{x}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right) \left( \sum_{k,t} \hat{\mathbf{z}}_{k,t} \hat{\mathbf{z}}_{k,t}^\top \right)^{-1} \quad (21)$$

$$\mathbf{S}^* = \frac{1}{KT} \left( \sum_{k,t} \mathbf{x}_{k,t} \mathbf{x}_{k,t}^\top - \mathbf{D}^* \sum_{k,t} \hat{\mathbf{z}}_{k,t} \mathbf{x}_{k,t}^\top \right) \quad (25)$$

These parameters define the observation model for neural data and will be used for decoding hand velocity in subsequent sections.

### 3.3 Using Kalman Filtering to Predict the Hand Velocity

Using the model parameters obtained in Section 3.2, we implemented a Kalman filter to compute the filtered posterior  $p(\mathbf{z}_{k',t} | \mathbf{x}_{k',0:t})$  for each trial  $k'$  in the test set. The mean of this posterior,  $\hat{\mathbf{z}}_{k',t}$ , serves as the best estimate of the latent state at time  $t$ . The hand velocity is then predicted as:

$$\hat{\mathbf{v}}_{k',t} = \mathbf{C} \hat{\mathbf{z}}_{k',t}. \quad (26)$$

The predicted velocities were submitted to the server, yielding an  $R^2$  score of 0.78.

**Why Are These Predictions Better?** The Kalman filter outperforms the baseline decoder (Section 2) for several reasons:

- **Temporal Dynamics:** Unlike the baseline method, the Kalman filter incorporates the sequential dependencies of the latent states through the state transition model. This allows it to leverage past observations to make more accurate predictions of the current state.
- **Noise Reduction:** The Kalman gain dynamically balances the contribution of the prior estimate and the new observation, effectively filtering out noise in the neural data.
- **Latent State Representation:** By utilizing the latent state dynamics modeled in the LDS, the Kalman filter captures the underlying neural representation of hand movement, leading to more precise predictions.

**Suitability for Online Decoding** The Kalman filter is highly suitable for online brain-computer interface (BCI) applications due to its recursive, computationally efficient nature, requiring only the current state estimate and the latest observation at each time step. The adaptive Kalman gain ensures robust performance in noisy environments.

However, its assumptions of static model parameters ( $\mathbf{A}$ ,  $\mathbf{C}$ ,  $\mathbf{Q}$ ,  $\mathbf{R}$ ) may limit adaptability to neural tuning changes over time, and delays in signal acquisition could impact real-time performance.

**Further Improvements** To enhance decoding performance, the following approaches are proposed:

- **ReFIT-Kalman Filter:** Adjusts decoder parameters during closed-loop operation based on user feedback, improving accuracy and robustness to neural activity changes. However, it requires additional calibration time.
- **Nonlinear State-Space Models:** Methods such as Recurrent Neural Networks (RNNs) and Gaussian Process Dynamical Models (GPDs) better capture complex neural dynamics but increase computational requirements.
- **Adaptive Noise Models:** Dynamically adjusting  $\mathbf{Q}$  and  $\mathbf{R}$  improves noise handling, though it adds computational overhead.
- **Cortical Dynamics Integration:** Incorporating models that explicitly account for cortical interactions can improve temporal decoding accuracy at the cost of higher model complexity.

## A Code Implementation

```
1  # [markdown]
2  # # 4G10 Coursework 2: predicting hand kinematics from neural data
3  #
4  #
5  # Please read carefully the last section of this notebook, which gives some of
6  # our expectations regarding your report.
7  #
8  # In this handout,
9  # - text that is underlined corresponds to things you have to do /
10 # implement.
11 #
12 # - text in bold corresponds to questions you need to answer in some form in
13 # your report.
14 #
15 #
16 #
17 #
18 # [markdown]
19 # # 1. Setup
20 #
21 # In this piece of 4G10 coursework, you will use neural data recorded in the
22 # primary motor cortex (M1) of a reaching monkey to predict the kinematics of
23 # the monkey's hand.
24 #
25 # The monkey initiated each trial by placing their hand in the center of a fronto-
26 # parallel screen. A target then appeared on the screen. The monkey had to wait
27 # for a 'go' cue before making a reaching movement towards the instructed
28 # target. The targets were placed in various positions in a virtual maze, which
29 # changed in each trial, forcing the monkey to make a variety of reaching
30 # movements across trials.
31 #
32 # The activity of  $N=162$  motor cortical neurons was recorded simultaneously,
33 # alongside the kinematics of the animal's hand.
34 #
35 # In the dataset presented below, all time series are partitioned into trials.
36 # Each trial begins at the go cue and lasts 800ms ( $T = 16$  bins of 50ms duration)
37 # roughly the duration of a reach.
38 #
39 #
40 # grab the data from the server
41 r = requests.get('http://4G10.cbl-cambridge.org/data.npz', stream = True)
42 data = np.load(BytesIO(r.raw.read()))
43 print(list(data.keys()))
44 #
45 # [markdown]
46 # Among other things (detailed later), this dictionary numerical arrays indexed
47 # by the following keys:
48 # - "hand_train" ( $2 \times 400 \times T$ ): 2D velocity (X/Y) of the monkey's hand in 400 '
49 # train' trials;
50 # - "neural_train" ( $N \times 400 \times T$ ): neural activity (spike counts) in the same 400
51 # 'train' trials;
52 # - "neural_test" ( $N \times 100 \times T$ ): neural activity (spike counts) in 100 'test'
53 # trials.
54 #
55 # E.g.:
56 #
57 #
```

```

45 # Load data
46 hand_train = data["hand_train"]
47 neural_train = data["neural_train"]
48 neural_test = data["neural_test"]
49
50 print(hand_train.shape)
51 print(neural_train.shape)
52 print(neural_test.shape)
53
54 # np.save('train_file.npy', hand_train[:,0:100,:].astype("float64"))
55
56 #
57 import numpy as np
58 import matplotlib.pyplot as plt
59
60 X = neural_train
61
62 # Compute the mean along dimension 1 (axis=1)
63 mean_dim1 = np.mean(X, axis=(1))
64
65 # Plot X over the mean of dim 1
66 plt.figure(figsize=(10, 6))
67 for i in range(5):
68     plt.plot(mean_dim1[i,:], linewidth=2, label='Mean (dim 1)') # Plot the mean
69 plt.title("Neural Data: Each Neuron vs. Mean (dim 1)")
70 plt.xlabel("Sample Index")
71 plt.ylabel("Value")
72 plt.legend(loc="upper right")
73 plt.grid(True)
74 plt.show()
75
76 # [markdown]
77 # The goal of this CW is to implement some of the modelling / decoding techniques
    you have been taught in lectures, to predict the monkey's 2D hand velocity in
    the 100 test trials for which you are only given neural activity. Your
    predictions will be based on the training data provided (hand_train,
    neural_train).
78
79 # [markdown]
80 # # 2. Baseline decoder: simple Gaussian smoothing + linear regression
81 #
82 # To establish a meaningful baseline, you will first implement a very simple two-
    stage decoder.
83 # In the first stage, you will smooth the spike count time series of each neuron
    by convolving it with a Gaussian filter of width  $\sigma$ ; in continuous time,
    such a Gaussian filter is given by  $f(t) \propto \exp(-t^2/2\sigma^2)$ .
84 # In the second stage, you will use ridge regression to learn an instantaneous
    linear decoder given by
85 #
86 # 
$$\hat{v}_{k,t} = W \tilde{x}_{k,t}$$

87 #
88 # where  $\hat{v}_{k,t} \in \mathbb{R}^2$  is the predicted velocity of the hand
    in test trial  $k$  and time bin  $t$ ,  $\tilde{x}_{k,t} \in \mathbb{R}^N$  is the
     $t$ -th time bin of the temporally smoothed spike counts in test trial
     $k$ , and  $W$  is a  $2 \times N$  matrix of decoding weights. Note that the hand
    velocity data has been centered already, so there is no need to include a bias
    term in the regression.
89 #
90 # The optimal ridge regression weights are given by
91 # 
$$W^* = V \tilde{X}^{\top} (\tilde{X} \tilde{X}^{\top} + \lambda I_N)^{-1}$$

92 # where  $V$  is the  $2 \times (400 \times 16)$  matrix of hand velocities from the training set
    (with all trials and time bins concatenated horizontally), and similarly  $\tilde{X}$ 

```



```

    tilde{X}$ is the $N \times (400 \times 16)$ matrix of smoothed neural spike counts in the
    training set.
93 #
94 # In the equation above, $\lambda$ is a regularisation parameter which helps
    protect against overfitting.
95 # The choice of value for this parameter is left up to you, so long as you can
    provide a justification (there are several sensible possibilities).
96 #
97 # The goal here is to make the best possible predictions you can of the held out
    hand velocity data in test trials, based on the neural activity in the same
    trials. When you are ready to test your predictions, you can submit them as a
    3D numpy array of shape 2 x 100 x 16 to http://4G10.cbl-cambridge.org (note:
    http, not https). If you get a HTTP error 400 back, it probably means the
    format is wrong. Your numpy array must be saved using the np.save("filename.
    npy", my_array) function; the server also expects the array to be of float64
    numerical type â this should be the default in numpy, but if in doubt you
    can always cast using my_array.astype("float64"). When you submit, please
    indicate your candidate number and choose "Simple Gaussian smoothing" in the
    dropdown list. Upon uploading, you will receive immediate feedback in the form
    of an  $R^2$  coefficient. The closer to 1, the better!
98 #
99 # - <u>Implement Gaussian temporal smoothing + ridge regression as outlined above
    </u>.
100 # - **How does the quality of hand velocity predictions vary with the smoothing
    window length  $\sigma$ ? How do you interpret that?** You might want to
    experiment with values between 20 and 80 ms.
101 # - **Comment on the suitability of this simple decoding strategy for online (on
    the fly) decoding of movement in a BMI context (consider e.g. feasibility,
    computational tractability, and accuracy). Can you think of a small
    modification to the above approach that would improve applicability to online
    decoding?** (bonus points for implementing it!)
102 # - The hand velocity data provided in data["hand_train"] had actually been
    shifted backward by 120ms relative to the neural data (and similarly for the
    test set, which was not given to you). **Can you speculate about why we did
    that**?
103
104 #
105 from scipy.ndimage import gaussian_filter1d
106
107 # Ridge regression
108 def ridge_regression(X, V, lambda):
109     N = X.shape[0]
110     XXt = X @ X.T
111     reg_matrix = lambda * np.eye(N)
112     W = V @ X.T @ np.linalg.inv(XXt + reg_matrix)
113     return W
114
115 # Prepare data
116 sigma = 1.6 # Set Gaussian smoothing window length (20ms, sigma=0.4) (80ms,
    sigma=1.6)
117 lambda_reg = 75 # Regularization parameter
118
119 # Smooth training neural data
120 smoothed_neural_train = gaussian_filter1d(neural_train, sigma=sigma, axis=2)
121
122 # Reshape training data for regression
123 N, trials, T = smoothed_neural_train.shape
124 smoothed_neural_train_reshaped = smoothed_neural_train.reshape(N, trials * T)
125 hand_train_reshaped = hand_train.reshape(2, trials * T)
126
127 # Compute ridge regression weights
128 W_star = ridge_regression(smoothed_neural_train_reshaped, hand_train_reshaped,
    lambda_reg)

```

```

129
130 # Predict on test data
131 smoothed_neural_test = gaussian_filter1d(neural_test, sigma=sigma, axis=2)
132 smoothed_neural_test_reshaped = smoothed_neural_test.reshape(N, -1)
133
134 # Predictions for test trials
135 predicted_hand_test = W_star @ smoothed_neural_test_reshaped
136 predicted_hand_test = predicted_hand_test.reshape(2, 100, T)
137
138 # Save predictions
139 np.save(f'hand_test_s={sigma}_l={lambda_reg}.npy', predicted_hand_test.astype("
    float64"))
140
141 #
142 from scipy.ndimage import gaussian_filter1d
143
144 # Ridge regression
145 def ridge_regression(X, V, lambd):
146     N = X.shape[0]
147     XXt = X @ X.T
148     reg_matrix = lambd * np.eye(N)
149     W = V @ X.T @ np.linalg.inv(XXt + reg_matrix)
150     return W
151
152 # Gaussian smoothing function (incremental for online use)
153 def online_gaussian_smooth(data, sigma, t):
154     """Applies Gaussian smoothing up to time step t."""
155     # smoothed = gaussian_filter1d(data[:, :, :t+1], sigma=sigma, axis=2, mode='
        constant', cval=0)
156     smoothed = gaussian_filter1d(data[:, :, :t+1], sigma=sigma, axis=2)
157     return smoothed[:, :, t] # Return the smoothed value at time t
158
159 # Prepare 5ata
160 sigma = 1.7 # Set Gaussian smoothing window length (20ms, sigma=0.4) (80ms,
        sigma=1.6)
161 lambda_reg = 75 # Regularization parameter
162
163 # Reshape training data for regression
164 N, trials, T = smoothed_neural_train.shape
165 predicted_hand_test_online = np.zeros((2, 100, T)) # To store online predictions
166 predicted_hand_train_online = np.zeros((2, 400, T)) # To store online
        predictions
167 smoothed_neural_test_online = np.zeros((N, 100, T)) # To store online
        predictions
168 smoothed_neural_train_online = np.zeros((N, 400, T)) # To store online
        predictions
169
170 ## train online
171 # for t in range(T): # Loop over each time step
172 #     smoothed_neural_train_online[:, :, t] = online_gaussian_smooth(neural_train,
        sigma, t)
173 # smoothed_neural_train_online_reshaped = smoothed_neural_train_online.reshape(N,
        400 * T)
174 # hand_train_reshaped = hand_train.reshape(2, 400 * T)
175 # W_star = ridge_regression(smoothed_neural_train_online_reshaped,
        hand_train_reshaped, lambda_reg)
176
177 ## train offline
178 W_star = ridge_regression(smoothed_neural_train_reshaped, hand_train_reshaped,
        lambda_reg)
179
180 for t in range(T): # Loop over each time step
181     # Smooth data incrementally up to time t

```

```

182     smoothed_neural_test_online[:, :, t] = online_gaussian_smooth(neural_test,
183                             sigma, t)
184     predicted_hand_test_online[:, :, t] = W_star @ smoothed_neural_test_online
185        [:, :, t]
186
187 # Save online predictions
188 print()
189 # print(predicted_hand_test_online.shape)
190 np.save(f'hand_test_online_incremental_s={sigma}_l={lambda_reg}.npy',
191         predicted_hand_test_online.astype("float64"))
192
193 print("Online prediction with incremental smoothing completed and saved.")
194
195 #
196 import numpy as np
197 import matplotlib.pyplot as plt
198 from scipy.ndimage import gaussian_filter1d
199
200 # Define the data and smoothing function
201 def online_gaussian_smooth(data, sigma, t):
202     """Applies Gaussian smoothing up to time step t."""
203     smoothed = gaussian_filter1d(data[:t+1], sigma=sigma, axis=0)
204     return smoothed[t] # Return the smoothed value at time t
205
206 # Initialize data
207 sigma = 1.0 # Example sigma value
208 test_zeros = np.zeros((16))
209 test_zeros[6] = 1
210
211 # Apply offline Gaussian smoothing
212 offline_smoothed = gaussian_filter1d(test_zeros, sigma=sigma)
213
214 # Apply online Gaussian smoothing
215 test_zeros_online = np.zeros((16))
216 for t in range(16): # Loop over each time step
217     test_zeros_online[t] = online_gaussian_smooth(test_zeros, sigma, t)
218
219 # Adjust x-axis values (multiply by 50)
220 time_steps = np.arange(16) * 50
221
222 # Plot the results
223 plt.figure(figsize=(3, 4))
224 plt.plot(time_steps, offline_smoothed, label='offline smoothed')
225 plt.plot(time_steps, test_zeros_online, label='online filtered')
226 plt.xlabel('Time (ms)')
227 plt.ylabel('Smoothed/Filtered Value')
228 plt.legend()
229 # plt.title('Comparison of Online and Offline Gaussian Smoothing')
230 plt.show()
231
232 #
233 import matplotlib.pyplot as plt
234 import numpy as np
235
236 # Data
237 sigma_offline = [20.0, 25.0, 30.0, 40.0, 50.0, 55.0, 60.0, 65.0, 70.0, 75.0,
238                 80.0]
239 r2_offline = [0.345, 0.375, 0.401, 0.428, 0.441, 0.444, 0.445, 0.445, 0.443,
240              0.440, 0.437]
241
242 sigma_online_trained_offline = [20.0, 25.0, 30.0, 40.0, 45.0, 50.0, 55.0, 60.0,
243                                65.0, 70.0, 75.0, 80.0, 85.0]
244 r2_online_trained_offline = [0.145, 0.202, 0.248, 0.307, 0.328, 0.345, 0.358,

```

```

0.367, 0.378, 0.378, 0.380, 0.380, 0.378]
239
240 sigma_online_trained_online = [20.0, 25.0, 30.0, 40.0, 50.0, 55.0, 60.0, 65.0,
70.0, 75.0, 80.0, 90.0, 95.0, 100.0]
241 r2_online_trained_online = [0.337, 0.353, 0.368, 0.391, 0.409, 0.416, 0.421,
0.426, 0.430, 0.433, 0.434, 0.436, 0.436, 0.435]
242
243 # Plot
244 plt.figure(figsize=(6, 4))
245
246 plt.plot(sigma_offline, r2_offline, label="Offline", marker='o', linestyle='-',
linewidth=2)
247 plt.plot(sigma_online_trained_offline, r2_online_trained_offline, label="Online,
trained offline", marker='s', linestyle='--', linewidth=2)
248 plt.plot(sigma_online_trained_online, r2_online_trained_online, label="Online,
trained online", marker='^', linestyle='-.', linewidth=2)
249
250 # Labels and legend
251 # plt.title("R2 vs  $\sigma$  Values", fontsize=14)
252 plt.xlabel("Smoothing Window  $\sigma$  (ms)", fontsize=12)
253 plt.ylabel("R2 Score", fontsize=12)
254 plt.legend(loc="best", fontsize=10)
255 plt.grid(True)
256
257 # Show plot
258 plt.tight_layout()
259 plt.show()
260
261 # [markdown]
262 # # 3. Kalman filter-based decoding
263 #
264 # We now turn to a more sophisticated decoder based on a Kalman filter/smoother.
265 #
266 # ### 3.1 An autoregressive prior for hand kinematics
267 #
268 # A 10-dimensional linear latent dynamical system (LDS; cf lecture notes) was pre-
trained for you on the hand velocity data in the training set; specifically,
we consider the following generative model:
269 #
270 # $$
271 # (1) \quad z_{k, 0} \sim \mathcal{N}(\mu_0, \Sigma_0) \quad \backslash \backslash
272 # (2) \quad z_{k, t+1} = A z_{k, t} + \epsilon_{k, t+1} \quad \text{with } \epsilon_{k, t+1} \sim \mathcal{N}(0, Q) \quad \backslash \backslash
273 # (3) \quad v_{k, t} = C z_{k, t} + \eta_{k, t} \quad \text{with } \eta_{k, t} \sim \mathcal{N}(0, R)
274 # $$
275 #
276 # where  $z_{k, t} \in \mathbb{R}^{10}$  is the latent state in time bin  $t$  of
trial  $k$ , and  $v_{k, t} \in \mathbb{R}^2$  is the corresponding hand velocity.
277 #
278 # The parameters of this LDS can be found in the same data dictionary as above,
with the following keys:
279 # - "hand_KF_A" (10 x 10): state matrix  $A$ 
280 # - "hand_KF_C" (2 x 10): output matrix  $C$ 
281 # - "hand_KF_mu0" (10 x 1): initial prior mean  $\mu_0$ 
282 # - "hand_KF_Sigma0" (10 x 10): initial prior covariance  $\Sigma_0$ 
283 # - "hand_KF_Q" (10 x 10): process noise covariance matrix  $Q$ 
284 # - "hand_KF_R" (2 x 2): observation noise covariance matrix  $R$ 
285 #
286 # <u>Write your own Kalman smoother implementation and use it to compute the mean
 $\hat{z}_{k, 1:T}$  of the smoothing distribution  $p(z_{k, t} \mid v_{k, 1:T})$ ,
for each trial  $k$  in the training set.</u>
287 #

```

```

288
289 #
290 # Retrieve parameters
291 A = data["hand_KF_A"]
292 C = data["hand_KF_C"]
293 mu0 = data["hand_KF_mu0"]
294 Sigma0 = data["hand_KF_Sigma0"]
295 Q = data["hand_KF_Q"]
296 R = data["hand_KF_R"]
297
298 #
299 print(mu0.shape)
300
301 #
302 import numpy as np
303
304
305 def kalman_filter(x, A, C, mu0, Sigma0, Q, R):
306     """
307     Kalman filter for the given LDS.
308
309     Parameters:
310         x: Observations (T x 2)
311         A: State transition matrix (10 x 10)
312         C: Observation matrix (2 x 10)
313         mu0: Initial mean (10 x 1)
314         Sigma0: Initial covariance (10 x 10)
315         Q: Process noise covariance (10 x 10)
316         R: Observation noise covariance (2 x 2)
317
318     Returns:
319         filtered_means: Smoothed means (T x 10)
320         filtered_covariances: Smoothed covariances (T x 10 x 10)
321     """
322     T = x.shape[0]
323     d = mu0.shape[0] # Dimensionality of latent state
324
325     # Initialize arrays
326     filtered_means = np.zeros((T, d))
327     filtered_covariances = np.zeros((T, d, d))
328
329     # Forward pass: Kalman filter
330     filtered_means[0] = mu0.ravel()
331     filtered_covariances[0] = Sigma0
332
333     for t in range(1, T, 1):
334         Pt = A @ filtered_covariances[t-1] @ A.T + Q
335         Kt = Pt @ C.T @ np.linalg.solve(C @ Pt @ C.T + R, np.eye(R.shape[0]))
336         filtered_means[t] = A @ filtered_means[t-1] + Kt @ (x[t] - C @ A @
337             filtered_means[t-1])
338         filtered_covariances[t] = Pt - Kt @ C @ Pt
339
340     return filtered_means, filtered_covariances
341
342 def kalman_smoother(x, A, C, mu0, Sigma0, Q, R):
343     """
344     Kalman smoother for the given LDS.
345
346     Parameters:
347         x: Observations (T x 2)
348         A: State transition matrix (10 x 10)
349         C: Observation matrix (2 x 10)
350         mu0: Initial mean (10 x 1)

```

```

350     Sigma0: Initial covariance (10 x 10)
351     Q: Process noise covariance (10 x 10)
352     R: Observation noise covariance (2 x 2)
353
354 Returns:
355     smoothed_means: Smoothed means (T x 10)
356     smoothed_covariances: Smoothed covariances (T x 10 x 10)
357 """
358 T = x.shape[0]
359 d = mu0.shape[0] # Dimensionality of latent state
360
361 # Forward pass: Kalman filter
362 filtered_means, filtered_covariances = kalman_filter(x, A, C, mu0, Sigma0, Q,
363     R)
364
365 # Backward pass: Kalman Smoother
366 smoothed_means = np.zeros((T, d))
367 smoothed_covariances = np.zeros((T, d, d))
368
369 smoothed_means[-1] = filtered_means[-1]
370 smoothed_covariances[-1] = filtered_covariances[-1]
371
372 for t in range(T-2, -1, -1):
373     Pt = A @ filtered_covariances[t] @ A.T + Q
374     Gt = filtered_covariances[t] @ A.T @ np.linalg.solve(Pt, np.eye(Pt.shape
375         [0]))
376     smoothed_means[t] = filtered_means[t] + Gt @ (smoothed_means[t+1] - A @
377         filtered_means[t])
378     smoothed_covariances[t] = filtered_covariances[t] + Gt @ (
379         smoothed_covariances[t+1] - Pt) @ Gt.T
380
381 return smoothed_means, smoothed_covariances
382
383 def kalman_all_trials(X, A, C, mu0, Sigma0, Q, R, mode='filter'):
384     # Velocity data shape: (N x 400 x 16)
385     # Output: (400 x 16 x 10)
386     d = A.shape[0]
387     n_trials = X.shape[1]
388     n_timesteps = X.shape[2]
389
390     # Initialize a container for smoothed means
391     all_kalmaned_means = np.zeros((d, n_trials, n_timesteps)) # (Latent Dim,
392         Trials, Timesteps)
393     all_kalmaned_covs = np.zeros((d, d, n_trials, n_timesteps)) # (Latent Dim,
394         Latent Dim, Trials, Timesteps)
395
396     # Loop over each trial
397     if mode=='filter':
398         for trial_idx in range(n_trials):
399             x_trial = X[:, trial_idx, :].T # Shape becomes (16 x 2)
400             kalmaned_means, kalmaned_covs = kalman_filter(x_trial, A, C, mu0,
401                 Sigma0, Q, R)
402             all_kalmaned_means[:, trial_idx, :] = kalmaned_means.T
403             all_kalmaned_covs[:, :, trial_idx, :] = np.transpose(kalmaned_covs,
404                 axes=(1,2,0))
405     elif mode=='smoother':
406         for trial_idx in range(n_trials):
407             x_trial = X[:, trial_idx, :].T # Shape becomes (16 x 2)
408             kalmaned_means, kalmaned_covs = kalman_smoother(x_trial, A, C, mu0,
409                 Sigma0, Q, R)
410             all_kalmaned_means[:, trial_idx, :] = kalmaned_means.T
411             all_kalmaned_covs[:, :, trial_idx, :] = np.transpose(kalmaned_covs,
412                 axes=(1,2,0))

```

```

403     else:
404         error
405     return all_kalmaned_means, all_kalmaned_covs
406
407 print(hand_train.shape)
408 smoothed_latent_means, _ = kalman_all_trials(hand_train, A, C, mu0, Sigma0, Q, R,
409     mode='smoother')
410 # smoothed_latent_means, _ = kalman_all_trials(hand_train, A, C, mu0, Sigma0, Q,
411     R, mode='filter')
412 print(smoothed_latent_means.shape)
413
414 # [markdown]
415 # ### 3.2 Building an LDS model of neural data using supervised learning
416 #
417 # Conceptually, the latents  $z_{\{k, 1:T\}}$  introduced above contain signals related
418 # to the velocity of the hand, its acceleration, and potentially higher-order
419 # derivatives too  $\hat{a}$  all signals which we have good reasons to suspect that
420 # neural activity in M1 is strongly related to. Eqs (1) and (2) above provide a
421 # good autoregressive prior model for the temporal dynamics of these signals,
422 # and you are now going to use this prior in a generative LDS model of _neural
423 # data_, substituting the hand-related likelihood (Eq 3) with a neural
424 # likelihood:
425 #
426 # 
$$x_{\{k, t\}} = D z_{\{k, t\}} + \xi_{\{k, t\}} \quad \text{with } \xi_{\{k, t\}} \sim \mathcal{N}(0, S)$$

427 #
428 # where  $x_{\{k, t\}}$  denotes neural spike counts in the  $t^{\text{th}}$  time bin of
429 # trial  $k$ .
430 #
431 # The combination of Eqs (1), (2) and (4) forms an LDS model which you will be
432 # able to invert using Kalman filtering to obtain a filtered posterior  $p(z_{\{k',$ 
433 #  $t\} | x_{\{k', 0:t\}})$  for any test trial  $k'$ . From there, you will use Eq. (3)
434 # to obtain a filtered predictive distribution for the hand velocity in each
435 # test trial,  $p(v_{\{k', t\} | x_{\{k', 0:t\}})$ .
436 #
437 # - Begin by centering the neural data (both training and testing sets) by
438 # removing, for each neuron, its mean activity across both trials and time in
439 # the training set.
440 #
441 # - Fit the likelihood parameters  $D$  and  $S$  through supervised learning, by
442 # maximizing the joint log-likelihood  $\log p(\hat{z}_{\{k, 1:T\}}, x_{\{k, 1:T\}})$ 
443 # averaged over all trials in the training set, where  $\hat{z}_{\{k, 1:T\}}$  is the
444 # posterior mean you obtained in Section 3.1.
445 #
446 # To do this, write down the average joint log likelihood and show that it is
447 # maximized by the following parameter settings:
448 #
449 # 
$$D^* = \left( \sum_{\{k, t\}} x_{\{k, t\}} \hat{z}_{\{k, t\}}^{\text{top}} \right) \left( \sum_{\{k, t\}} \hat{z}_{\{k, t\}} \hat{z}_{\{k, t\}}^{\text{top}} \right)^{-1}$$

450 #
451 # 
$$S^* = \frac{1}{KT} \left( \sum_{\{k, t\}} x_{\{k, t\}} x_{\{k, t\}}^{\text{top}} - D^* \sum_{\{k, t\}} \hat{z}_{\{k, t\}} x_{\{k, t\}}^{\text{top}} \right)$$

452 # where  $K$  is the number of trials in the training set. Include your derivations in
453 # your report.
454 #
455 import numpy as np
456
457 def center_neural_data(neural_train, neural_test):
458     """
459     Centers neural data by subtracting the mean activity of each neuron across
460     all trials and time bins in the training set.
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999

```



```

440 Parameters:
441     neural_train: Training neural data ( $N \times K \times T$ )
442     neural_test: Testing neural data ( $N \times K' \times T$ )
443
444 Returns:
445     centered_train: Centered training data ( $N \times K \times T$ )
446     centered_test: Centered testing data ( $N \times K' \times T$ )
447     """
448     # Compute mean activity across trials and time bins
449     mean_activity = np.mean(neural_train, axis=(1, 2), keepdims=True)
450
451     # Subtract mean from both training and testing sets
452     centered_train = neural_train - mean_activity
453     centered_test = neural_test - mean_activity
454
455     return centered_train, centered_test
456
457 def fit_likelihood_params(X, smoothed_means):
458     """
459     Fits the parameters  $D$  and  $S$  using the provided formulas.
460
461     Parameters:
462         X: Centered training neural data ( $N \times K \times T$ )
463         smoothed_means: Posterior means of latent states  $z$  ( $d \times K \times T$ )
464     Returns:
465         D_star: Fitted observation matrix ( $N \times d$ )
466         S_star: Fitted observation noise covariance ( $N \times N$ )
467     """
468     # Extract dimensions
469     N, K, T = X.shape #  $N$  = neurons,  $K$  = trials,  $T$  = timesteps
470     d, _, _ = smoothed_means.shape #  $d$  = latent dimensionality
471
472     # Initialize accumulators with correct dimensions
473     sum_xzT = np.zeros((N, d)) # Shape ( $N \times d$ )
474     sum_zzT = np.zeros((d, d)) # Shape ( $d \times d$ )
475     sum_xxT = np.zeros((N, N)) # Shape ( $N \times N$ )
476
477     # Iterate over trials and time bins
478     for k in range(K):
479         for t in range(T):
480             # Extract data for time  $t$  and trial  $k$ 
481             x_kt = X[:, k, t].reshape(-1, 1) # Shape ( $N \times 1$ )
482             z_kt = smoothed_means[:, k, t].reshape(-1, 1) # Shape ( $d \times 1$ )
483
484             # Accumulate sums
485             sum_xzT += x_kt @ z_kt.T # Outer product ( $N \times 1$ ) @ ( $1 \times d$ ) -> ( $N \times d$ )
486             sum_zzT += z_kt @ z_kt.T # Outer product ( $d \times 1$ ) @ ( $1 \times d$ ) -> ( $d \times d$ )
487             sum_xxT += x_kt @ x_kt.T # Outer product ( $N \times 1$ ) @ ( $1 \times N$ ) -> ( $N \times N$ )
488
489     # Compute  $D_{\text{star}}$ 
490     D_star = sum_xzT @ np.linalg.solve(sum_zzT, np.eye(d)) # Shape ( $N \times d$ )
491     # Compute  $S_{\text{star}}$ 
492     S_star = (sum_xxT - D_star @ sum_xzT.T) / (K * T) # Shape ( $N \times N$ )
493
494     return D_star, S_star
495
496
497 centered_neural_train, centered_neural_test = center_neural_data(neural_train,
498     neural_test)

```



```

499 # Fit the likelihood parameters
500 D_star, S_star = fit_likelihood_params(centered_neural_train,
    smoothed_latent_means)
501 # plt.figure()
502 # plt.imshow(D_star)
503 # plt.figure()
504 # plt.imshow(S_star)
505
506
507 # [markdown]
508 # ### 3.3 Using Kalman filtering to predict the hand velocity
509 #
510 # - Based on the model obtained in Section 3.2, <u>write your own Kalman filter
    implementation and compute the filtered posterior  $p(z_{k'}, t | x_{k'}, 0:t)$ 
    $ for each trial $k'$ in the test set.</u> What we are really interested in is
    the mean  $\bar{z}_{k', t}$  of this filtered posterior, as our best prediction
    of the momentary hand velocity  $v_{k', t}$  is then given by  $\hat{v}_{k', t} = C \bar{z}_{k', t}$ .
511 # - <u>Submit your predictions to http://4G10.cbl-cambridge.org (note: http, not
    https) in the same format as described in Section 2. Please select "Kalman
    filtering" in the dropdown list. Once again, you will receive immediate
    feedback in the form of an  $R^2$  coefficient.</u> **Include this result in
    your report, and discuss; in particular, why do you think these predictions
    are much better than those of Section 2?***
512 # - As in Section 2, **comment on the suitability of this more sophisticated
    decoding strategy for online decoding of movement in a BCI context**.
513 # - Finally, **what approach(es) would you suggest to improve decoding
    performance even further? Include a discussion of the tradeoffs that would
    arise with these alternative approaches.** (max 1 page in your report)
514
515 #
516 # filtered_latent_means, _ = kalman_all_trials(centered_neural_train, A, D_star,
    mu0, Sigma0, Q, S_star, mode='filter')
517 filtered_latent_means, _ = kalman_all_trials(centered_neural_test, A, D_star, mu0
    , Sigma0, Q, S_star, mode='filter')
518 v_pred = np.einsum('ij,jkl->ikl', C, filtered_latent_means)
519
520 # np.save(f'hand_test_online_incremental_s={sigma}_l={lambda_reg}.npz',
    predicted_hand_test_online.astype("float64"))
521 np.save(f'hand_test_kalman.npz', v_pred.astype("float64"))
522
523
524
525 #
526 filtered_latent_means, _ = kalman_all_trials(centered_neural_train, A, D_star,
    mu0, Sigma0, Q, S_star, mode='filter')
527
528 idx1 = 4
529 idx2 = 60
530 plt.plot(smoothed_latent_means[idx1,idx2,:], label='from v')
531 plt.plot(filtered_latent_means[idx1,idx2,:], label='from x')
532 plt.legend()
533
534 # [markdown]
535 # # Writing up
536 #
537 #
538 # Please write up your findings in a report to be submitted on Moodle in PDF
    format, and **include all your code in the Appendix**. Please clearly include
    your candidate number, NOT your name, on the front page. Your report should
    address all the questions raised in this notebook, be structured around the
    Sections of this notebook, and **be a maximum of five A4 pages** excluding any
    Appendix (minimum font size 11pt, minimum margins 1.5cm on each side).

```

539 #  
540 # *You are very much encouraged to think of data/results visualisations to best support the exposition of your results. You are also encouraged to report on any specific problems/difficulties that arose in your implementation of the various algorithms, and how you addressed those.*